

Careers, Overcompliance, and Coups

A Global Game of Career Pressure in Authoritarian Security Apparatuses

Kristoff (Yifan) Cao*

August 31, 2026

Abstract

Why do some officers in authoritarian security services volunteer for the regime’s dirtiest work while others of apparently identical rank and background plot against it? Gläsel and Scharpf (2026) argue that both behaviors are products of the same cause: blocked careers. I formalize this claim as a global game played by a continuum of security officers who differ in career prospects and hold private information about regime strength. Blocked officers choose between *detouring*—escalating volunteered brutality as a bid for out-of-turn promotion—and *forcing*—joining a coup. The two responses are strategic mirror images: detouring is a bet on regime survival, forcing a bet on collapse, and detourers directly suppress forcers. Three results follow. First, the equilibrium is a monotone partition of the signal space in which forcing occurs at low signals and detouring at high signals, with ordinary compliance occupying an interior band that vanishes when blockage is severe; the model therefore *derives* rather than assumes the bimodality that motivates it. Second, the critical state at which the regime falls admits a closed form that is exact for every noise level and is independent of the precision of officers’ information—a sharp null prediction that separates this account from information-manipulation theories of regime change. Third, and centrally, regime survival is non-monotone in career blockage: the critical state is a strictly convex quadratic in the mass of blocked officers, so moderate blockage strictly *expands* the set of fundamentals under which the regime survives while severe blockage destroys it. A dictator who chooses blockage optimally sets it at an interior level b^{**} , increasing in the coup-suppressing value of detourers and decreasing in the anticipated severity of post-transition punishment of collaborators. Blocked careers are simultaneously the regime’s principal vulnerability and its principal defense.

Keywords: global games; coups; authoritarian regimes; repression; bureaucratic careers; coup-proofing.

*Preliminary and incomplete. Comments welcome. This paper formalizes the argument of Gläsel and Scharpf (2026). All errors are my own.

1 Introduction

Authoritarian security services routinely produce two kinds of extreme behavior. Some officers volunteer, unprompted, for assignments that others avoid: interrogations, disappearances, mass killings. Others conspire against the ruler they serve. These behaviors are usually explained separately—the first by ideology, selection, or situational pressure; the second by ambition, grievance, or opportunity. [Gläsel and Scharpf \(2026\)](#) argue that they share a cause. Officers whose careers are blocked, who cannot expect promotion through ordinary channels, respond in one of two extreme ways: they *detour*, volunteering for the most burdensome and psychologically straining tasks in order to be noticed, or they *force*, attempting to seize by violence what the promotion ladder will not give them. The empirical counterpart of this argument is striking: in autocratic Argentina, secret police organizations were dominated not by the ablest officers but by underachievers ([Scharpf and Gläsel, 2020](#)).

This paper asks what follows once that argument is taken seriously as a strategic claim. Detouring and forcing are not two independent reactions to a common stimulus. They interact. An officer who volunteers for the regime’s dirty work is, among other things, an officer who will help crush a coup; an officer who joins a coup succeeds only if enough others do, and only if the regime’s enforcers are too few to stop them. The same career blockage that manufactures the threat also manufactures the defense. Whether blockage on net stabilizes or destabilizes a regime is therefore not obvious, and cannot be settled by intuition.

I model the security apparatus as a global game ([Carlsson and van Damme, 1993](#); [Morris and Shin, 1998](#)). A continuum of officers observes noisy private signals of the regime’s underlying strength θ and chooses among normal compliance (N), detouring (T), and forcing (C). A coup succeeds when the mass of forcers, net of the mass of detourers weighted by their suppressive effectiveness α , exceeds θ . Officers differ in one respect that matters: a fraction b of them are career-blocked, so that the normal ladder is worth nothing to them, while the remainder are on track and comply. The parameter b is the model’s measure of career pressure, and in the final section it becomes the dictator’s choice variable.

Three features of the environment do the work. First, detouring pays only if the regime survives: promotions are regime-specific, and officers who made their names doing the regime’s dirty work are the first targets of post-transition punishment ([Nalepa, 2010](#)). Detouring is therefore a bet on survival, exactly as forcing is a bet on collapse, and the two are ordered in the signal space in opposite directions. Second, detouring is congestible:

the more officers compete for the same handful of out-of-turn promotions, the further each must go to stand out. The escalation cost ϕx_T is what makes competitive brutality self-limiting, and it turns out to be the binding constraint on how much career pressure a ruler can profitably impose. Third, detourers protect the regime, at rate α , against the forcers drawn from the same population.

Results. Under a single-crossing condition on the escalation cost, the equilibrium is a monotone partition: officers force at low signals, comply at intermediate signals, and detour at high signals (Proposition 1). The compliance band is empty—the apparatus is *polarized*, with every blocked officer either detouring or plotting—exactly when $(R - \chi)/(R + P) \geq (\kappa + L + \phi b/2)/(\delta + L)$ (Proposition 3). This is the model’s formal statement of the Gläsel–Scharpf bimodality: it is derived from primitives rather than posited, and it holds precisely when both extremes are individually attractive relative to a career that is going nowhere.

The critical state has the closed form

$$\theta^* = b[(1 + \alpha)q - \alpha], \quad q = \frac{R + L - \chi + \kappa + \phi b/2}{R + P + L + \delta},$$

where q is the share of blocked officers who force at the critical state. Two things are worth noticing. The expression contains no noise parameter: θ^* is exact for every $\sigma > 0$, not merely in the vanishing-noise limit (Proposition 2). Improving officers’ information changes *who* plots but not *when* the regime falls. And θ^* is a strictly convex quadratic in b with $\theta^*(0) = 0$, so career blockage is beneficial to the regime at the margin whenever $q < \alpha/(1 + \alpha)$: the marginal blocked officer is more likely to become a shield than a threat (Proposition 5).

The dictator’s problem follows immediately. Survival is maximized at an interior level of blockage

$$b^{**} = \frac{1}{\phi} \left[\frac{\alpha}{1 + \alpha} (R + P + L + \delta) - (R + L - \chi + \kappa) \right],$$

and blockage beyond $2b^{**}$ leaves the regime strictly more fragile than an apparatus with open careers (Proposition 7). Coup-proofing through career blockage obeys a Laffer curve. Its peak rises with the suppressive value of detourers α and with the punishment of failed coups P , and falls with the coup prize R , the psychological cost of dirty work κ , the scarcity of promotion slots ϕ , and—a result with contemporary bite—the anticipated severity L of post-transition justice. Credible transitional-justice threats do not merely deter individual collaboration; they shrink the amount of career pressure a ruler can

profitably impose, because they degrade the shield that pressure was meant to produce.

Contribution. Existing formal treatments of coups and regime change are, with few exceptions, binary-action coordination games: agents attack or they do not (Bueno de Mesquita, 2010; Shadmehr and Bernhardt, 2011; Edmond, 2013; Casper and Tyson, 2014; Little, 2017; Tyson and Smith, 2018). The substantive claim in Gläsel and Scharpf (2026) cannot be represented in that space, because it is a claim about which of *two* deviations a frustrated officer takes, and about how those two deviations interact. The model here is a three-action global game in which one action is a strategic complement (forcing) and the other is congestible and suppressive (detouring). That combination is not supermodular, so the standard selection results (Frankel et al., 2003) do not apply off the shelf; Section 7 shows that the failure is not merely technical. If detouring is placed in the interior of the signal space—the natural first guess—the aggregate pressure function becomes non-monotone and the game admits multiple critical states, with disjoint regions of collapse separated by a survival window. The ordering of actions in the signal space is thus not a modeling convenience but a substantive commitment with consequences for whether the model has a determinate prediction at all.

The paper also connects two literatures that have developed apart. The loyalty–competence literature (Egorov and Sonin, 2011; Zakharov, 2016) asks whom a dictator should appoint; the coup-proofing literature (Quinlivan, 1999; Powell, 2012; Svolik, 2012) asks how he should structure the forces once appointed. Career blockage is an instrument that operates on both margins at once, and its optimum is interior for a reason neither literature isolates: the same act that produces loyal enforcers produces the men who would replace him.

Section 2 situates the paper. Section 3 sets up the game and discusses its assumptions. Section 4 characterizes equilibrium. Section 5 derives the survival results. Section 6 solves the dictator’s problem. Section 7 presents the multiplicity result. Section 8 states testable implications, and Section 9 discusses extensions and limitations. Proofs are in Appendix A.

2 Related literature

Global games of regime change. The technology is standard (Carlsson and van Damme, 1993; Morris and Shin, 1998, 2003). Political applications have studied revolutionary entrepreneurs (Bueno de Mesquita, 2010), punishment dilemmas and public sig-

nals (Shadmehr and Bernhardt, 2011), state information manipulation (Edmond, 2013), learning across time (Angeletos et al., 2007), and the interaction of mass protest with elite coordination (Casper and Tyson, 2014; Tyson and Smith, 2018; Little, 2017). All treat the coordinating agents as choosing between two actions. The present model adds a third action that is neither participation nor abstention but a distinct form of *over*-participation, and shows that the third action changes the comparative statics qualitatively rather than incrementally.

Loyalty, competence, and coup-proofing. Egorov and Sonin (2011) and Zakharov (2016) derive the dictator’s trade-off between able and loyal subordinates. Detouring is a close cousin of loyalty: it is a costly display whose value lies in what it signals about the officer’s dependence on the regime. What is new here is that the display is *produced* by the promotion structure rather than being an exogenous type, so the dictator’s instrument acts on the supply of loyalty directly. The coup-proofing literature (Quinlivan, 1999; Powell, 2012; Svolik, 2012; Geddes et al., 2018) treats counterbalancing and parallel forces as costly diversions of capacity. Career blockage is different in kind: it is free in resource terms and its cost is entirely strategic, which is why its optimum is interior rather than a corner.

Repression and its agents. Dragu and Przeworski (2019) model the moral hazard inherent in delegating preventive repression. Gläsel and Scharpf (2026) and Scharpf and Gläsel (2020) supply the microhistorical foundation for the present paper: security organizations staffed disproportionately by officers with poor prospects, whose volunteered brutality is best read as career behavior. Nalepa (2010) documents the post-transition exposure of collaborators that gives the parameter L its empirical content.

3 Model

3.1 Players, types, and actions

There is a dictator D and a unit continuum of security officers indexed by $i \in [0, 1]$. Section 6 gives D a choice; until then he is passive.

Each officer has a *career type* $\tau_i \in \{H, L\}$, privately known. A fraction $b \in [0, 1]$ are *blocked* ($\tau_i = L$): the ordinary promotion ladder is closed to them. The remaining $1 - b$ are *on track* ($\tau_i = H$). The parameter b is the model’s index of *career pressure*: it is

the share of the officer corps for whom normal service has no career value. Types are independent of everything else.

Each officer chooses $a_i \in \{N, T, C\}$:

- N (*normal compliance*): serve, obey routine orders, wait for promotion.
- T (*detouring*): volunteer for the regime’s most burdensome and psychologically straining assignments as a bid for out-of-turn advancement.
- C (*forcing*): join a coup attempt.

3.2 The state, information, and regime collapse

The latent state $\theta \in \mathbb{R}$ is the regime’s *strength*: the cohesion, resources, and residual loyalty that a coup must overcome. Following standard practice, θ has an improper uniform prior on $\mathbb{R}$; this is the limit of a diffuse proper prior and rules out common knowledge of fundamentals (Morris and Shin, 2003). Officer i observes

$$s_i = \theta + \sigma \varepsilon_i, \quad \varepsilon_i \stackrel{iid}{\sim} F,$$

with $\sigma > 0$ and F continuous, strictly increasing, with density $f > 0$ on $\mathbb{R}$, symmetric about zero and single-peaked. *High s_i is a signal of regime strength; low s_i is a signal of weakness.*

Write the aggregate action shares

$$x_T(\theta) = \int_0^1 \mathbf{1}\{a_i = T\} di, \quad x_C(\theta) = \int_0^1 \mathbf{1}\{a_i = C\} di.$$

The coup succeeds—the regime *falls*—if and only if

$$x_C - \alpha x_T \geq \theta, \quad \alpha \in (0, 1). \tag{1}$$

The parameter α is the rate at which detourers suppress forcers. Two readings are available and the model does not distinguish them: detourers physically defend the regime, or—more interestingly—detourers, being the officers most invested in the regime’s survival, are the ones who denounce plotters, so α measures the informational density of the internal security net.

Remark 1 (Orientation). Condition (1) makes θ regime *strength*: a larger θ is a higher hurdle. Consequently a high private signal makes an officer *less* willing to force, and the

forcing region of the signal space lies to the *left*. This is the opposite of the orientation one obtains by reading θ as “weakness” while retaining (1), and Section 7 shows the two are not interchangeable.

3.3 Payoffs

Careers are regime-specific: a promotion is worth having only if the regime that granted it survives. Let w_τ denote the value of the normal ladder to a type- τ officer, with the normalization $w_L = 0$. Payoffs are

$$U_i(N) = w_{\tau_i} \cdot \mathbf{1}\{\text{survive}\}, \quad (2)$$

$$U_i(T) = \delta \cdot \mathbf{1}\{\text{survive}\} - L \cdot \mathbf{1}\{\text{fall}\} - \kappa - \phi x_T, \quad (3)$$

$$U_i(C) = R \cdot \mathbf{1}\{\text{fall}\} - P \cdot \mathbf{1}\{\text{survive}\} - \chi_{\tau_i}, \quad (4)$$

with $\delta, \kappa, \phi, R, P, \chi_\tau, L > 0$. The interpretation of each term:

- δ is the value of the out-of-turn promotion that successful detouring buys. It is collected only if the regime survives.
- $L \geq 0$ is the punishment a detourer suffers if the regime falls. Officers who distinguished themselves by volunteered brutality are the most exposed to post-transition prosecution, purge, or private retribution (Nalepa, 2010). L is the parameter through which anticipated transitional justice enters.
- κ is the up-front personal cost of the dirty work itself.
- ϕx_T is the *escalation cost*. Promotion slots reserved for demonstrated loyalty are scarce; the more rivals volunteer, the further each must go to be noticed. This term is paid regardless of the regime’s fate, because the work is done before the outcome is known. It is the model’s representation of competitive brutality, and it is what prevents detouring from being a free lunch.
- R , P , and χ_τ are the coup prize, the punishment for a failed attempt, and the fixed cost of conspiring. The latter is type-dependent: an officer with a career has more to lose from exposure, so $\chi_H > \chi_L$. Write $\chi \equiv \chi_L$.

Assumption 1 (On-track officers comply). $w_H \geq \delta$ and $\chi_H \geq R$.

Under Assumption 1, N strictly dominates both T and C for type- H officers (Lemma 4 in the appendix), so only blocked officers act. The assumption imports one half of the Gläsel–Scharpf claim—that officers with prospects do not take extreme career actions—and leaves the other half, the split of blocked officers between the two extremes, to be derived. Section 9 discusses relaxing it.

For a blocked officer, writing $\Pi(s) \equiv \Pr\{\text{fall} \mid s\}$ and using $w_L = 0$:

$$\mathbb{E}U(N \mid s) = 0, \tag{5}$$

$$\mathbb{E}U(T \mid s) = \delta(1 - \Pi(s)) - L\Pi(s) - \kappa - \phi\mathbb{E}[x_T \mid s], \tag{6}$$

$$\mathbb{E}U(C \mid s) = (R + P)\Pi(s) - P - \chi. \tag{7}$$

3.4 Discussion of the specification

Three modeling choices deserve comment.

Why detouring must be state-contingent. If the return to detouring did not depend on the regime’s fate, $\mathbb{E}U(T \mid s)$ would vary with s only through the congestion term, and at signals high enough that nobody detours, detouring would again become attractive; no compliance region could survive at the top of the signal space. Making the promotion prize contingent on survival is the minimal fix, and it is also the correct one: it is what makes detouring the mirror image of forcing rather than an unrelated third option. The lustration term L then completes the symmetry—forcers are punished if the coup fails, detourers if it succeeds.

Why congestion is a cost rather than a rationing probability. Modeling the promotion contest explicitly, with μ reserved slots rationed among x_T volunteers, would give a survival-contingent prize $\delta \min\{1, \mu/x_T\}$. This is tractable and yields the same qualitative results, but the up-front escalation cost ϕx_T is a better description of the phenomenon: competition among volunteers manifests as an arms race in demonstrated ruthlessness, which is borne whether or not the promotion arrives. The linear form is a first-order approximation to any decreasing per-volunteer return; nothing below depends on linearity except the closed forms.

Why b rather than a continuum of career types. A continuous distribution of prospects w would deliver type-dependent cutoffs and an additional margin of comparative statics, at the cost of closed forms. The binary specification isolates the quantity of interest—the *mass* of officers for whom the ladder is worthless—and makes it the dictator’s instrument.

Table 1: Notation

Symbol	Meaning
θ	latent regime strength
$s_i = \theta + \sigma \varepsilon_i$	officer i 's private signal; $\varepsilon_i \sim F$
b	mass of career-blocked officers (career pressure)
x_T, x_C	mass detouring, mass forcing
α	rate at which detourers suppress forcers
δ	value of out-of-turn promotion (paid on survival)
L	post-transition punishment of detourers (paid on collapse)
κ	up-front personal cost of dirty work
ϕ	escalation cost per unit of detouring rivals
R, P, χ	coup prize, punishment for failure, cost of conspiring
$\Pi(s)$	posterior probability of collapse given s
θ^*	critical state: regime falls iff $\theta \leq \theta^*$
q_C, q_T, q	equilibrium quantiles defined in Proposition 1

4 Equilibrium

4.1 Aggregates and the critical state

Fix a monotone strategy for blocked officers, described by cutoffs $\hat{s}_C \leq \hat{s}_T$ with $a_i = C$ for $s_i < \hat{s}_C$, $a_i = N$ for $s_i \in [\hat{s}_C, \hat{s}_T)$, and $a_i = T$ for $s_i \geq \hat{s}_T$. Then

$$x_C(\theta) = b F\left(\frac{\hat{s}_C - \theta}{\sigma}\right), \quad x_T(\theta) = b \left[1 - F\left(\frac{\hat{s}_T - \theta}{\sigma}\right)\right]. \quad (8)$$

Define the *net pressure function*

$$\Psi(\theta) \equiv x_C(\theta) - \alpha x_T(\theta) - \theta. \quad (9)$$

The regime falls iff $\Psi(\theta) \geq 0$.

Lemma 1 (Unique critical state). *Under (8), $\Psi'(\theta) < -1 < 0$ for all θ , and $\lim_{\theta \rightarrow -\infty} \Psi = +\infty$, $\lim_{\theta \rightarrow +\infty} \Psi = -\infty$. Hence there is a unique θ^* with $\Psi(\theta^*) = 0$, and the regime falls if and only if $\theta \leq \theta^*$.*

The content of Lemma 1 is that x_C is decreasing and x_T is *increasing* in θ , so both the direct effect of strength and the induced composition effect push in the same direction. Section 7 shows what happens when they do not.

Given Lemma 1, the posterior probability of collapse is

$$\Pi(s) = \Pr\{\theta \leq \theta^* \mid s\} = 1 - F\left(\frac{s - \theta^*}{\sigma}\right), \quad (10)$$

strictly decreasing in s .

4.2 Monotone best responses

Assumption 2 (Moderate escalation). $\phi b \sup_x (f \star f)(x) < (\delta + L) \inf_{y \in \mathcal{Y}} f(y)$, where $f \star f$ is the density of $\varepsilon_1 + \varepsilon_2$ and $\mathcal{Y}$ is any compact interval containing the standardized cutoffs $(\hat{s}_C - \theta^*)/\sigma$ and $(\hat{s}_T - \theta^*)/\sigma$.

Assumption 2 says that the escalation externality is not so strong that it reverses the state-dependence of the detouring payoff. It is satisfied for ϕ small relative to $\delta + L$, and it is the only place where an inequality on ϕ is needed.

Lemma 2 (Monotone partition). *Fix any θ^* and any monotone profile of others. Then $\mathbb{E}U(C \mid s)$ is strictly decreasing in s , and under Assumption 2 $\mathbb{E}U(T \mid s)$ is strictly increasing in s . Consequently a blocked officer's best response is a monotone partition of the form*

$$a_i = \begin{cases} C, & s_i < \hat{s}_C, \\ N, & \hat{s}_C \leq s_i < \hat{s}_T, \\ T, & s_i \geq \hat{s}_T, \end{cases}$$

with the middle region possibly empty.

The ordering is the substantive heart of the model. Forcing is a bet on collapse and is taken when collapse looks likely; detouring is a bet on survival and is taken when survival looks likely; compliance is what remains for officers who are unsure. The two extremes bracket ordinary behavior rather than lying on the same side of it.

4.3 The Laplacian property

The step that delivers closed forms is the following exact result, which requires no limit.

Lemma 3 (Beliefs at the detouring cutoff). *Under the improper uniform prior and (8), an officer whose signal equals the detouring cutoff believes the mass of detourers is distributed*

as $b(1 - U)$ with $U \sim \text{Uniform}[0, 1]$. Hence

$$\mathbb{E}[x_T \mid s = \hat{s}_T] = \frac{b}{2} \quad \text{for every } \sigma > 0.$$

Proof sketch. At $s = \hat{s}_T$ the posterior is $\theta = \hat{s}_T - \sigma\varepsilon$ with $\varepsilon \sim F$, so $x_T = b[1 - F(\varepsilon)]$; and $F(\varepsilon) \sim U[0, 1]$. $\square$

4.4 Characterization

Proposition 1 (Equilibrium). *Define*

$$q_C \equiv \frac{R - \chi}{R + P}, \quad q_T \equiv \frac{\kappa + L + \phi b/2}{\delta + L}, \quad q \equiv \frac{R + L - \chi + \kappa + \phi b/2}{R + P + L + \delta},$$

and assume $R > \chi$ and $\delta > \kappa + \phi b/2$ so that $q_C, q_T \in (0, 1)$. Under Assumptions 1–2 the game has a monotone equilibrium, described by exactly one of two regimes.

(i) **Guarded regime** ($q_C < q_T$). The compliance band is non-empty. The cutoffs and the critical state are

$$\hat{s}_C = \theta^* + \sigma F^{-1}(q_C), \quad \hat{s}_T = \theta^* + \sigma F^{-1}(q_T), \quad \boxed{\theta^* = b[q_C - \alpha(1 - q_T)]}.$$

(ii) **Polarized regime** ($q_C \geq q_T$). No blocked officer complies. There is a single cutoff $\hat{s}$ and

$$\hat{s} = \theta^* + \sigma F^{-1}(q), \quad \boxed{\theta^* = b[(1 + \alpha)q - \alpha]}.$$

In this case q is exactly the share of blocked officers who force when $\theta = \theta^*$, and $1 - q$ the share who detour.

The economics of the three quantiles is transparent. q_C is set by the coup gamble alone: an officer forces when the odds of success exceed the odds ratio implied by R , P , and χ . q_T is set by the career gamble alone: an officer detours when the odds of survival are good enough to justify κ , the escalation cost, and the risk of ending up on a lustration list. q is the head-to-head comparison and lies between them.

Proposition 2 (Uniqueness and noise-irrelevance). *The monotone equilibrium of Proposition 1 is unique. Moreover θ^* does not depend on σ : the same critical state obtains for every noise level, not merely in the limit $\sigma \rightarrow 0$. The cutoffs themselves satisfy $\hat{s} - \theta^* = O(\sigma)$ and converge to θ^* as $\sigma \rightarrow 0$.*

Remark 2. Proposition 2 claims uniqueness within monotone strategies. Extending it to all strategies by iterated dominance, as in the binary case, would require the game to be supermodular. It is not: forcing is a strategic complement, while detouring is a strategic substitute through both the escalation cost and the suppression term, so Frankel et al. (2003) does not apply directly. Establishing—or refuting—global uniqueness for this class is the paper’s main open technical question.

4.5 When the middle disappears

Proposition 3 (Polarization). *The compliance band is empty if and only if*

$$\frac{R - \chi}{R + P} \geq \frac{\kappa + L + \phi b/2}{\delta + L}. \quad (11)$$

When (11) holds, $x_T(\theta) + x_C(\theta) = b$ at every θ : career blockage translates one-for-one into extreme behavior, and officers’ private signals determine only how the blocked population divides between the two extremes.

Condition (11) is the formal counterpart of the Gläsel–Scharpf bimodality. It holds when coups are attractive relative to their cost (R high, χ low), when the loyalty route is richly rewarded and not too personally costly (δ high, κ low), and when lustration is not severe (L low). Note that (11) becomes *harder* to satisfy as b grows: congestion among volunteers eventually restores a compliance band. Career blockage polarizes the apparatus, but crowding at the loyalty gate limits how far.

Proposition 4 (Bimodal response to blockage). *Evaluated at the critical state, an increase in career blockage raises the mass of forcers, $\partial x_C(\theta^*)/\partial b = q + b q' > 0$; and it raises the mass of detourers, $\partial x_T(\theta^*)/\partial b = (1 - q) - b q' > 0$, whenever $b q' < 1 - q$, where $q' = \frac{\phi/2}{R+P+L+\delta}$. Both extremes therefore expand together, while ordinary compliance contracts.*

Figure 1 illustrates the two regimes.

5 Career pressure and regime survival

The regime survives on (θ^*, ∞) and falls on $(-\infty, \theta^*]$. For any proper prior over θ with full support, the survival probability is strictly decreasing in θ^* , so θ^* is a sufficient statistic for regime fragility and I work with it directly.




Figure 1: Expected payoffs of a blocked officer. Forcing is decreasing in the signal, detouring increasing, compliance flat. In panel (a) the two extremes cross above zero and no officer complies; in panel (b) they cross below zero and a compliance band opens between the cutoffs. Escalation costs are held at their equilibrium value $\phi b/2$. Panel (a): $R=P=1$, $\chi=0.2$, $\delta=1.5$, $L=0.2$, $\kappa=0.05$. Panel (b): $R=P=1$, $\chi=0.6$, $\delta=0.9$, $L=0.4$, $\kappa=0.15$. Both: $\phi=0.8$, $b=0.5$, $\alpha=0.6$, $\sigma=0.3$, logistic noise.

Proposition 5 (The critical state is a convex quadratic in blockage). *In the polarized regime,*

$$\theta^*(b) = Ab + Bb^2, \quad A = (1 + \alpha)q_0 - \alpha, \quad B = \frac{(1 + \alpha)\phi}{2D} > 0, \quad (12)$$

where $D \equiv R + P + L + \delta$ and $q_0 \equiv (R + L - \chi + \kappa)/D = q|_{b=0}$. In the guarded regime the same form holds with $A = q_C - \alpha + \alpha \frac{\kappa+L}{\delta+L}$ and $B = \frac{\alpha\phi}{2(\delta+L)} > 0$. In both cases $\theta^*(0) = 0$ and θ^* is strictly convex in b .

Corollary 1 (Coup-proofing Laffer curve). *Suppose $A < 0$, equivalently (polarized regime) $q_0 < \alpha/(1 + \alpha)$. Then:*

- (i) $\theta^*(b) < 0 = \theta^*(0)$ for all $b \in (0, -A/B)$: moderate career blockage strictly expands the set of fundamentals under which the regime survives;
- (ii) θ^* is minimized at $b^{**} = -A/(2B)$, with $\theta^*(b^{**}) = -A^2/(4B) < 0$;
- (iii) $\theta^*(b) > 0$ for all $b > 2b^{**}$: blockage beyond twice the optimum leaves the regime strictly more fragile than an apparatus with open careers.

If instead $A \geq 0$, θ^* is strictly increasing on $[0, 1]$ and any career blockage is destabilizing.




Figure 2: The coup-proofing Laffer curve. Lower θ^* means the regime survives a wider range of fundamentals. Parameters as in Figure 1(a) with $\alpha = 0.6$, giving $b^{**} = 0.422$ and break-even blockage $2b^{**} = 0.844$.

The threshold $q_0 < \alpha/(1 + \alpha)$ has a clean reading. q_0 is the fraction of blocked officers who force at the critical state in an uncongested apparatus. The marginal blocked officer becomes a forcer with probability q_0 and a detourer with probability $1 - q_0$; the first raises pressure by one unit, the second lowers it by α . The net effect is $q_0 - \alpha(1 - q_0) = (1 + \alpha)q_0 - \alpha$, which is A . Blockage is stabilizing at the margin precisely when the shield the marginal officer provides outweighs the threat he poses.

The convexity—the reason the effect must eventually reverse—comes entirely from congestion. As b rises, more officers crowd the loyalty gate, each must escalate further, and detouring becomes less attractive relative to forcing: $q'(b) > 0$. Without congestion ($\phi = 0$) we would have $B = 0$ and blockage would be either uniformly good or uniformly bad for the regime. *The scarcity of promotion slots is what makes the dictator's problem interior.*

Proposition 6 (Information precision is irrelevant). $\partial\theta^*/\partial\sigma = 0$. *Changes in the precision of officers' private information shift the cutoffs but leave the critical state, and hence the regime's survival set, unchanged.*

This is a genuinely restrictive prediction and it distinguishes the model sharply from theories in which the ruler's central instrument is control of information (Edmond, 2013; Guriev and Treisman, 2020). It relies on the diffuse prior: with public signals or a proper prior the neutrality breaks, and—as in Morris and Shin (2004)—sufficiently precise public information can restore multiplicity. Section 9 returns to this.

Table 2: Comparative statics of optimal career blockage b^{**}

Parameter	$\partial b^{**}/\partial \cdot$	Interpretation
α suppressive value of detourers	+	The denser the internal informer net, the more the ruler gains from manufacturing detourers, and the more careers he should block.
ϕ scarcity of promotion slots	−	Crowding at the loyalty gate degrades the shield; top-heavy apparatuses can afford less blockage.
P punishment of failed coups	+	Harsher retribution suppresses forcing, freeing the ruler to press careers harder.
R coup prize	−	Where seizing power pays well, blockage produces plotters faster than shields.
χ cost of conspiring	+	Surveillance and compartmentalization are complements to career pressure.
δ value of out-of-turn promotion	+	Richer loyalty rewards tilt the blocked toward detouring.
κ personal cost of dirty work	−	When the work is more repellent, fewer of the blocked take the loyalty route.
L post-transition punishment of collaborators	−	Credible transitional justice degrades the shield and <i>lowers</i> the optimal amount of career pressure.

6 The dictator's problem

Let the dictator choose $b \in [0, 1]$ before θ is realized, to maximize the probability of survival, equivalently to minimize θ^* . Career blockage is resource-free: it is a property of the promotion schedule, not an allocation of forces. Its entire cost is strategic.

Proposition 7 (Optimal career blockage). *In the polarized regime, if $q_0 < \alpha/(1 + \alpha)$ and $b^{**} \leq 1$, the dictator's optimum is interior and equals*

$$b^{**} = \frac{1}{\phi} \left[\frac{\alpha}{1 + \alpha} (R + P + L + \delta) - (R + L - \chi + \kappa) \right], \quad (13)$$

delivering $\theta^*(b^{**}) = -\frac{[\alpha - (1 + \alpha)q_0]^2}{4(1 + \alpha)\phi/(2D)} < 0$. If $q_0 \geq \alpha/(1 + \alpha)$ the optimum is $b^{**} = 0$; if the bracket in (13) exceeds ϕ , it is $b^{**} = 1$.

Two entries in Table 2 are worth dwelling on.

Surveillance and career pressure are complements. b^{**} is increasing in both χ and P .

A regime that has invested in detecting and punishing conspiracy can afford to block more careers, because the marginal blocked officer is then more likely to detour than to plot. This rationalizes an empirical regularity that the coup-proofing literature treats as two separate strategies: heavy internal-security investment and notoriously congested promotion ladders tend to appear together.

Transitional justice has an equilibrium effect on repression's supply. b^{**} is decreasing in L . The standard argument for credible post-transition accountability is deterrence at the individual level: an officer who expects to be prosecuted is less willing to commit atrocities. The model adds a second, structural channel. Because lustration risk makes detouring less attractive, it degrades the very shield that career blockage was meant to produce, and a forward-looking ruler responds by blocking fewer careers. The prediction is that credible international or domestic accountability regimes should be associated not only with less volunteered brutality but with *looser* promotion bottlenecks inside security services.

7 Why the ordering of actions matters

Lemma 2 places detouring at the top of the signal space. It is tempting to place it in the middle instead—compliance for the most sanguine officers, detouring under intermediate pressure, forcing for the most desperate—which is what one obtains if the detouring payoff is specified without state contingency. This subsection shows that the alternative is not a harmless relabeling: it destroys the uniqueness of the critical state.

Suppose blocked officers played C on $(-\infty, \hat{s}_C)$, T on $[\hat{s}_C, \hat{s}_T)$, and N on $[\hat{s}_T, \infty)$. Then $x_T(\theta) = b \left[F\left(\frac{\hat{s}_T - \theta}{\sigma}\right) - F\left(\frac{\hat{s}_C - \theta}{\sigma}\right) \right]$ is *hump-shaped* in θ , and

$$\Psi'(\theta) = -\frac{(1 + \alpha)b}{\sigma} f\left(\frac{\hat{s}_C - \theta}{\sigma}\right) + \frac{\alpha b}{\sigma} f\left(\frac{\hat{s}_T - \theta}{\sigma}\right) - 1, \quad (14)$$

whose sign is indeterminate.

Proposition 8 (Multiplicity under interior detouring). *Under the interior-detouring ordering, if the detouring band is wide relative to the noise, $\hat{s}_T - \hat{s}_C \gg \sigma$, then $\Psi'(\theta) > 0$ for θ near $\hat{s}_T$ whenever*

$$\sigma < \alpha b f(0).$$

In that case Ψ is non-monotone and may cross zero three times, so the regime has two disjoint collapse regions separated by an interval of survival.




Figure 3: Net pressure $\Psi(\theta) = x_C - \alpha x_T - \theta$ under the two orderings. $b = 1$, $\alpha = 0.8$, $\sigma = 0.05$, $\hat{s}_C = -0.5$, $\hat{s}_T = -0.2$, logistic noise. The regime collapses where $\Psi \geq 0$. With detouring in an interior band the shield melts as θ rises and Ψ turns upward, so a second, anomalous collapse region (shaded) opens at fundamentals *stronger* than those the regime survives.

The mechanism is worth stating plainly. When detourers occupy an interior band, an increase in regime strength eventually pushes the marginal officer *out* of detouring and into compliance. The shield melts away faster than the fundamental improves, and the regime becomes more vulnerable as it becomes stronger. The resulting survival window—an interval of intermediate fundamentals on which the regime survives, flanked on both sides by collapse—is a substantively interesting object, but it also means the model has no determinate prediction about when regimes fall.

Figure 3 contrasts the two orderings. Numerically, with $b = 1$, $\alpha = 0.8$, $\sigma = 0.05$, logistic noise, $\hat{s}_C = -0.5$ and $\hat{s}_T = -0.2$, the interior-detouring specification yields three crossings of Ψ (at $\theta \approx -0.431$, -0.107 , -0.022), while the specification of Lemma 2 yields exactly one. Since $f(0) = 1/4$ for the standardized logistic, the threshold in Proposition 8 is $\sigma < 0.2$, and indeed raising σ to 0.3 restores monotonicity.

8 Empirical implications

The model yields several implications that are, in principle, testable with the personnel-record data that motivates it.

1. **Composition.** Both volunteered brutality and coup participation are concentrated among career-blocked officers; on-track officers do neither. This is the premise on

the H side and the finding of [Scharpf and Gläsel \(2020\)](#), and it is the model’s weakest empirical claim because it is assumed rather than derived.

2. **Separation by perceived regime strength.** Among blocked officers, the ratio of detourers to forcers is increasing in perceived regime strength. Periods of apparent regime security should see surges of volunteered brutality; periods of apparent fragility should see conspiracies drawn from the same population. This is the model’s sharpest prediction and the one that most clearly separates it from a pure dispositional account, which would predict the two behaviors to covary positively rather than to trade off.
3. **Non-monotone survival in blockage.** Regime durability should be an inverted U in the severity of promotion bottlenecks, measurable as the ratio of eligible officers to available slots by cohort. Both very open and very blocked apparatuses should be more coup-prone than moderately blocked ones.
4. **Complementarity of surveillance and blockage.** Regimes with more effective internal-security detection (higher χ , P) and denser informer networks (higher α) should exhibit *more* career blockage in equilibrium, not less.
5. **Transitional justice and promotion structure.** Where post-transition accountability is credible (higher L), optimal blockage falls. Periods following regional precedents of prosecution should be associated with loosening promotion bottlenecks and declining rates of volunteered brutality.
6. **Escalation.** The responsiveness of volunteered brutality to career blockage is dampened where promotion slots are scarcest, because escalation costs rise with the number of rivals.
7. **An informational null.** By Proposition 6, variation in the precision of officers’ information about regime strength changes who plots but not the timing of collapse. Finding that better-informed officer corps produce systematically earlier or later collapses would falsify this specification and point toward models with public signals.

9 Extensions and limitations

Dynamics. The model is static, which is its most serious limitation given that careers are inherently dynamic. The natural extension is an overlapping-cohort promotion tour-

nament in which detouring today buys a probability of promotion tomorrow, so that b becomes an endogenous state variable evolving with the ruler’s past promotion choices. This would allow the model to speak to the observed life-cycle of authoritarian security services, in which blockage accumulates as cohorts age into a narrowing pyramid.

Endogenous α and denunciation. Treating α as a parameter suppresses an interesting margin: detourers suppress coups partly by informing, and the decision to inform is itself strategic. Endogenizing α as an equilibrium denunciation rate would connect the model to [Dragu and Przeworski \(2019\)](#).

A richer dictator. The ruler here chooses only b . Letting him also choose δ subject to a rent budget, or the precision of monitoring, would let the model speak to the composition of coup-proofing portfolios. Note from Table 2 that δ and b are complements, so a ruler with a large rent budget should exhibit both generous loyalty rewards and severe blockage—a testable joint prediction.

On-track officers. Assumption 1 rules out coups by well-placed officers, which is empirically too strong: many coups are led from the top. Relaxing it requires type-dependent cutoffs and loses the closed forms, but the qualitative structure should survive, with H officers holding a strictly lower forcing cutoff.

Public information. Adding a public signal breaks Proposition 6 and, if the public signal is precise enough, can restore multiplicity ([Morris and Shin, 2004](#)). Since authoritarian rulers actively manage public information about their own strength, this is the most substantively important extension.

Global uniqueness. As noted after Proposition 2, the game is not supermodular, and uniqueness is established here only within monotone strategies.

10 Conclusion

Blocked careers inside an authoritarian security apparatus produce two extreme responses, and the two are not independent. Officers who volunteer for the regime’s dirty work are betting on its survival; officers who plot are betting on its collapse; and the first group suppresses the second. Formalizing this as a three-action global game yields a determinate

equilibrium—forcing at low signals, detouring at high signals, compliance in between when it survives at all—and a closed form for the critical state that is exact at every noise level.

The central substantive result is that career pressure is neither straightforwardly stabilizing nor straightforwardly destabilizing. Because the marginal blocked officer becomes a shield with probability $1 - q$ and a threat with probability q , and because crowding at the loyalty gate raises q as blockage grows, the critical state is a convex quadratic in blockage. Moderate blockage makes a regime robust to fundamentals that would otherwise destroy it; severe blockage makes it more fragile than an apparatus with open careers. A ruler who understood this would block careers deliberately, and stop.

A Proofs

Lemma 4. *Under Assumption 1, N strictly dominates T and C for type- H officers.*

Proof. For a type- H officer, $\mathbb{E}U(N \mid s) = w_H(1 - \Pi(s))$. Comparing with detouring,

$$\mathbb{E}U(N \mid s) - \mathbb{E}U(T \mid s) = (w_H - \delta)(1 - \Pi(s)) + L\Pi(s) + \kappa + \phi\mathbb{E}[x_T \mid s] > 0$$

whenever $w_H \geq \delta$, since $L, \kappa, \phi \geq 0$ and the last three terms are non-negative with $\kappa > 0$. Comparing with forcing,

$$\mathbb{E}U(N \mid s) - \mathbb{E}U(C \mid s) = w_H(1 - \Pi) - [(R + P)\Pi - P - \chi_H] = w_H(1 - \Pi) + P(1 - \Pi) + \chi_H - R\Pi.$$

Since $\Pi \in [0, 1]$, the right-hand side is bounded below by its value at $\Pi = 1$, namely $\chi_H - R \geq 0$, with strict inequality for $\Pi < 1$. Hence N is weakly dominant and strictly so on the interior. $\square$

Proof of Lemma 1. From (8),

$$\frac{dx_C}{d\theta} = -\frac{b}{\sigma}f\left(\frac{\hat{s}_C - \theta}{\sigma}\right) \leq 0, \quad \frac{dx_T}{d\theta} = \frac{b}{\sigma}f\left(\frac{\hat{s}_T - \theta}{\sigma}\right) \geq 0,$$

so $\Psi'(\theta) = -\frac{b}{\sigma}f\left(\frac{\hat{s}_C - \theta}{\sigma}\right) - \frac{ab}{\sigma}f\left(\frac{\hat{s}_T - \theta}{\sigma}\right) - 1 \leq -1 < 0$. Since $x_C \in [0, b]$ and $x_T \in [0, b]$ are bounded, $\Psi(\theta) \rightarrow \pm\infty$ as $\theta \rightarrow \mp\infty$. Strict monotonicity and continuity give a unique zero. $\square$

Proof of Lemma 2. By (10), $\Pi'(s) = -\frac{1}{\sigma}f\left(\frac{s - \theta^*}{\sigma}\right) < 0$. Then $\frac{d}{ds}\mathbb{E}U(C \mid s) = (R + P)\Pi'(s) < 0$. For detouring, rewrite (6) as $\mathbb{E}U(T \mid s) = \delta - (\delta + L)\Pi(s) - \kappa - \phi\mathbb{E}[x_T \mid s]$, so

$$\frac{d}{ds}\mathbb{E}U(T \mid s) = \frac{1}{\sigma} \left[(\delta + L)f\left(\frac{s - \theta^*}{\sigma}\right) - \phi b(f \star f)\left(\frac{\hat{s}_T - s}{\sigma}\right) \right],$$

using $\mathbb{E}[x_T \mid s] = b\mathbb{E}_\varepsilon[1 - F(\frac{\hat{s}_T - s}{\sigma} + \varepsilon)]$ and differentiating under the integral. Assumption 2 makes the bracket strictly positive on the relevant compact interval. Since $\mathbb{E}U(N) = 0$ is constant, $\mathbb{E}U(C \mid \cdot)$ strictly decreasing and $\mathbb{E}U(T \mid \cdot)$ strictly increasing, each pairwise difference is strictly monotone and crosses zero at most once, which yields the stated partition. $\square$

Proof of Lemma 3. Conditional on $s = \hat{s}_T$ and the improper uniform prior, $\theta = \hat{s}_T - \sigma\varepsilon$

with $\varepsilon \sim F$. Substituting into (8),

$$x_T = b \left[1 - F\left(\frac{\hat{s}_T - \theta}{\sigma}\right) \right] = b[1 - F(\varepsilon)].$$

Since F is continuous and strictly increasing, $F(\varepsilon) \sim U[0, 1]$, so $x_T \sim b(1 - U)$ and $\mathbb{E}[x_T \mid \hat{s}_T] = b/2$. No approximation is used, so the result holds for every $\sigma > 0$. $\square$

Proof of Proposition 1. Guarded regime. At $\hat{s}_C$ a blocked officer is indifferent between C and N : $(R+P)\Pi(\hat{s}_C) - P - \chi = 0$, so $\Pi(\hat{s}_C) = \frac{P+\chi}{R+P}$ and, by (10), $F\left(\frac{\hat{s}_C - \theta^*}{\sigma}\right) = 1 - \Pi(\hat{s}_C) = \frac{R-\chi}{R+P} = q_C$. At $\hat{s}_T$ he is indifferent between N and T : using Lemma 3,

$$\delta(1 - \Pi(\hat{s}_T)) - L\Pi(\hat{s}_T) - \kappa - \frac{\phi b}{2} = 0 \implies F\left(\frac{\hat{s}_T - \theta^*}{\sigma}\right) = 1 - \Pi(\hat{s}_T) = \frac{\kappa + L + \phi b/2}{\delta + L} = q_T.$$

Evaluating (8) at $\theta = \theta^*$ gives $x_C(\theta^*) = b q_C$ and $x_T(\theta^*) = b(1 - q_T)$. Substituting into $\Psi(\theta^*) = 0$ yields $\theta^* = b[q_C - \alpha(1 - q_T)]$. The cutoff expressions follow by inverting F .

Polarized regime. If $q_C \geq q_T$ the two zero-crossings are ordered so that at every s at least one of C, T delivers a non-negative payoff, and N is never a strict best response. Indifference between C and T at the single cutoff $\hat{s}$ gives

$$(R + P)\Pi - P - \chi = \delta(1 - \Pi) - L\Pi - \kappa - \frac{\phi b}{2},$$

so $\Pi(R + P + \delta + L) = P + \chi + \delta - \kappa - \phi b/2$ and $F\left(\frac{\hat{s} - \theta^*}{\sigma}\right) = 1 - \Pi = q$. Then $x_C(\theta^*) = bq$, $x_T(\theta^*) = b(1 - q)$, and $\Psi(\theta^*) = 0$ gives $\theta^* = b[(1 + \alpha)q - \alpha]$. $\square$

Proof of Proposition 2. In each regime the indifference conditions determine the standardized cutoffs $(\hat{s} - \theta^*)/\sigma$ uniquely and *independently of* θ^* , because $\mathbb{E}[x_T \mid \hat{s}] = b/2$ by Lemma 3 regardless of where the cutoff lies, and Π depends on the signal only through the standardized distance to θ^* . Given those quantiles, $\Psi(\theta^*) = 0$ determines θ^* uniquely by Lemma 1. Neither step involves σ , so θ^* is invariant to σ ; the cutoffs themselves are $\theta^* + \sigma F^{-1}(\cdot)$ and collapse to θ^* as $\sigma \rightarrow 0$. $\square$

Proof of Proposition 3. N is chosen on a non-degenerate interval iff $\hat{s}_C < \hat{s}_T$, iff $q_C < q_T$ by monotonicity of F^{-1} , which is the negation of (11). When (11) holds, every blocked officer plays C or T , so $x_C + x_T = b$ pointwise in θ . $\square$

Proof of Proposition 5. Polarized regime: substitute $q = q_0 + q'b$ with $q' = \frac{\phi/2}{\delta}$ into $\theta^* = b[(1 + \alpha)q - \alpha]$ to get $\theta^* = b[(1 + \alpha)q_0 - \alpha] + (1 + \alpha)q'b^2$, which is (12) with $B =$

$(1+\alpha)\phi/(2D) > 0$. Guarded regime: substitute $q_T = \frac{\kappa+L}{\delta+L} + \frac{\phi b/2}{\delta+L}$ into $\theta^* = b[q_C - \alpha(1-q_T)]$ and collect terms. $\square$

Proof of Corollary 1. $\theta^*(b) = Ab + Bb^2$ with $B > 0$ is a strictly convex parabola through the origin. Its roots are 0 and $-A/B$; with $A < 0$ the second root is positive and $\theta^* < 0$ strictly between them. The vertex is at $b = -A/(2B)$ with value $-A^2/(4B)$. For $b > -A/B = 2b^{**}$, $\theta^* > 0 > \theta^*(0)$. $\square$

Proof of Proposition 7. Minimizing (12) gives $b^{**} = -A/(2B)$ with $A = (1+\alpha)q_0 - \alpha$ and $B = (1+\alpha)\phi/(2D)$, so $b^{**} = \frac{\alpha-(1+\alpha)q_0}{(1+\alpha)\phi/D} = \frac{D}{\phi} \left[\frac{\alpha}{1+\alpha} - q_0 \right]$. Substituting $Dq_0 = R + L - \chi + \kappa$ gives (13). The corner cases follow from convexity. $\square$

Proof of Proposition 8. Under the interior-detouring ordering, $x_T(\theta) = b \left[F\left(\frac{\hat{s}_T - \theta}{\sigma}\right) - F\left(\frac{\hat{s}_C - \theta}{\sigma}\right) \right]$, and differentiating $\Psi = x_C - \alpha x_T - \theta$ gives (14). Take $\theta = \hat{s}_T$. If $\hat{s}_T - \hat{s}_C \gg \sigma$ then $f\left(\frac{\hat{s}_C - \hat{s}_T}{\sigma}\right) \approx 0$ while $f\left(\frac{\hat{s}_T - \theta}{\sigma}\right) = f(0)$, so $\Psi'(\hat{s}_T) \approx \frac{\alpha b}{\sigma} f(0) - 1$, which is strictly positive iff $\sigma < \alpha b f(0)$. Since $\Psi \rightarrow +\infty$ as $\theta \rightarrow -\infty$ and $\Psi \rightarrow -\infty$ as $\theta \rightarrow +\infty$, an interior interval of increase permits three crossings; whether three occur depends on whether the local maximum of Ψ to the right of the first crossing is non-negative. $\square$

References

- Angeletos, G.-M., Hellwig, C., and Pavan, A. (2007). Dynamic global games of regime change: Learning, multiplicity, and the timing of attacks. *Econometrica*, 75(3):711–756.
- Bueno de Mesquita, E. (2010). Regime change and revolutionary entrepreneurs. *American Political Science Review*, 104(3):446–466.
- Carlsson, H. and van Damme, E. (1993). Global games and equilibrium selection. *Econometrica*, 61(5):989–1018.
- Casper, B. A. and Tyson, S. A. (2014). Popular protest and elite coordination in a coup d’état. *Journal of Politics*, 76(2):548–564.
- Dragu, T. and Przeworski, A. (2019). Preventive repression: Two types of moral hazard. *American Political Science Review*, 113(1):77–87.
- Edmond, C. (2013). Information manipulation, coordination, and regime change. *Review of Economic Studies*, 80(4):1422–1458.
- Egorov, G. and Sonin, K. (2011). Dictators and their viziers: Endogenizing the loyalty–competence trade-off. *Journal of the European Economic Association*, 9(5):903–930.
- Frankel, D. M., Morris, S., and Pauzner, A. (2003). Equilibrium selection in global games with strategic complementarities. *Journal of Economic Theory*, 108(1):1–44.
- Geddes, B., Wright, J., and Frantz, E. (2018). *How Dictatorships Work: Power, Personalization, and Collapse*. Cambridge University Press, Cambridge.
- Gläsel, C. and Scharpf, A. (2026). *Making a Career in Dictatorship: The Secret Logic behind Repression and Coups*. Oxford University Press, Oxford.
- Guriev, S. and Treisman, D. (2020). A theory of informational autocracy. *Journal of Public Economics*, 186:104158.
- Little, A. T. (2017). Coordination, learning, and coups. *Journal of Conflict Resolution*, 61(1):204–234.
- Morris, S. and Shin, H. S. (1998). Unique equilibrium in a model of self-fulfilling currency attacks. *American Economic Review*, 88(3):587–597.

- Morris, S. and Shin, H. S. (2003). Global games: Theory and applications. In Dewatripont, M., Hansen, L. P., and Turnovsky, S. J., editors, *Advances in Economics and Econometrics: Theory and Applications, Eighth World Congress*, pages 56–114. Cambridge University Press, Cambridge.
- Morris, S. and Shin, H. S. (2004). Coordination risk and the price of debt. *European Economic Review*, 48(1):133–153.
- Nalepa, M. (2010). *Skeletons in the Closet: Transitional Justice in Post-Communist Europe*. Cambridge University Press, Cambridge.
- Powell, J. (2012). Determinants of the attempting and outcome of coups d’état. *Journal of Conflict Resolution*, 56(6):1017–1040.
- Quinlivan, J. T. (1999). Coup-proofing: Its practice and consequences in the Middle East. *International Security*, 24(2):131–165.
- Scharpf, A. and Gläsel, C. (2020). Why underachievers dominate secret police organizations: Evidence from autocratic Argentina. *American Journal of Political Science*, 64(4):791–806.
- Shadmehr, M. and Bernhardt, D. (2011). Collective action with uncertain payoffs: Coordination, public signals, and punishment dilemmas. *American Political Science Review*, 105(4):829–851.
- Svolik, M. W. (2012). *The Politics of Authoritarian Rule*. Cambridge University Press, Cambridge.
- Tyson, S. A. and Smith, A. (2018). Dual-layered coordination and political instability: Repression, co-optation, and the role of information. *Journal of Politics*, 80(1):44–58.
- Zakharov, A. V. (2016). The loyalty–competence trade-off in dictatorships and outside options for subordinates. *Journal of Politics*, 78(2):457–466.